

softmax

May 8, 2018

1 Softmax exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

This exercise is analogous to the SVM exercise. You will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient
- use a validation set to **tune the learning rate and regularization strength**
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
In [17]: import random
import numpy as np
from hwk5_2.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

from __future__ import print_function

# Reference: https://github.com/Halfish/cs231n/tree/master/assignment1

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

```

In [18]: def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000, num_dev=1000):
        """
        Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
        it for the linear classifier. These are the same steps as we used for the
        SVM, but condensed to a single function.
        """

        # Load the raw CIFAR-10 data
        cifar10_dir = 'hmk5_2/datasets/cifar-10-batches-py'

        X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

        # subsample the data
        mask = list(range(num_training, num_training + num_validation))
        X_val = X_train[mask]
        y_val = y_train[mask]
        mask = list(range(num_training))
        X_train = X_train[mask]
        y_train = y_train[mask]
        mask = list(range(num_test))
        X_test = X_test[mask]
        y_test = y_test[mask]
        mask = np.random.choice(num_training, num_dev, replace=False)
        X_dev = X_train[mask]
        y_dev = y_train[mask]

        # Preprocessing: reshape the image data into rows
        X_train = np.reshape(X_train, (X_train.shape[0], -1))
        X_val = np.reshape(X_val, (X_val.shape[0], -1))
        X_test = np.reshape(X_test, (X_test.shape[0], -1))
        X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

        # Normalize the data: subtract the mean image
        mean_image = np.mean(X_train, axis = 0)
        X_train -= mean_image
        X_val -= mean_image
        X_test -= mean_image
        X_dev -= mean_image

        # add bias dimension and transform into columns
        X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
        X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
        X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
        X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

        return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev

# Cleaning up variables to prevent loading data multiple times (which may cause memory

```

```

try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev = get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
print('dev data shape: ', X_dev.shape)
print('dev labels shape: ', y_dev.shape)

```

```

Clear previously loaded data.
Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)

```

1.1 Softmax Classifier

Your code for this section will all be written inside `cs231n/classifiers/softmax.py`.

```

In [19]: # First implement the naive softmax loss function with nested loops.
# Open the file cs231n/classifiers/softmax.py and implement the
# softmax_loss_naive function.

from hwk5_2.classifiers.softmax import softmax_loss_naive
import time

# Generate a random softmax weight matrix and use it to compute the loss.
W = np.random.randn(3073, 10) * 0.0001
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As a rough sanity check, our loss should be something close to -log(0.1).
print('loss: %f' % loss)
print('sanity check: %f' % (-np.log(0.1)))

```

```
loss: 2.411441
sanity check: 2.302585
```

1.2 Inline Question 1:

Why do we expect our loss to be close to $-\log(0.1)$? Explain briefly.**

Your answer: The probability for each class is 0.1 since we have 10 classes. Therefore, the log loss is $-\log(0.1)$.

```
In [20]: # Complete the implementation of softmax_loss_naive and implement a (naive)
# version of the gradient that uses nested loops.
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As we did for the SVM, use numeric gradient checking as a debugging tool.
# The numeric gradient should be close to the analytic gradient.
from hwk5_2.gradient_check import grad_check_sparse
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)

# similar to SVM case, do another gradient check with regularization
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

```
numerical: -0.814163 analytic: -0.814163, relative error: 2.837687e-09
numerical: 1.002279 analytic: 1.002279, relative error: 3.929047e-09
numerical: 0.594972 analytic: 0.594972, relative error: 1.953488e-08
numerical: 1.797474 analytic: 1.797474, relative error: 2.765331e-08
numerical: 0.063704 analytic: 0.063704, relative error: 1.292642e-07
numerical: 0.188682 analytic: 0.188682, relative error: 1.793670e-07
numerical: -2.539226 analytic: -2.539226, relative error: 1.491037e-09
numerical: 0.815371 analytic: 0.815371, relative error: 5.148228e-08
numerical: -0.741378 analytic: -0.741378, relative error: 6.687431e-09
numerical: 1.023604 analytic: 1.023604, relative error: 1.995192e-08
numerical: 3.163773 analytic: 3.163773, relative error: 1.823740e-08
numerical: -0.437112 analytic: -0.437112, relative error: 3.192093e-09
numerical: -5.682758 analytic: -5.682758, relative error: 5.466139e-09
numerical: -1.724527 analytic: -1.724527, relative error: 2.861011e-08
numerical: 2.520643 analytic: 2.520643, relative error: 4.503707e-09
numerical: -0.240242 analytic: -0.240242, relative error: 7.630763e-08
numerical: -1.677741 analytic: -1.677741, relative error: 1.413622e-08
numerical: -2.649825 analytic: -2.649825, relative error: 7.792913e-10
numerical: 1.273833 analytic: 1.273833, relative error: 6.748062e-09
numerical: -0.652263 analytic: -0.652263, relative error: 2.075119e-08
```

```
In [21]: # Now that we have a naive implementation of the softmax loss function and its gradient,
# implement a vectorized version in softmax_loss_vectorized.
```

```

# The two versions should compute the same results, but the vectorized version should
# much faster.
tic = time.time()
loss_naive, grad_naive = softmax_loss_naive(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('naive loss: %e computed in %fs' % (loss_naive, toc - tic))

from hwk5_2.classifiers.softmax import softmax_loss_vectorized
tic = time.time()
loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

# As we did for the SVM, we use the Frobenius norm to compare the two versions
# of the gradient.
grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
print('Loss difference: %f' % np.abs(loss_naive - loss_vectorized))
print('Gradient difference: %f' % grad_difference)

```

```

naive loss: 2.411441e+00 computed in 0.098400s
vectorized loss: 2.411441e+00 computed in 0.008782s
Loss difference: 0.000000
Gradient difference: 0.000000

```

```

In [22]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of over 0.35 on the validation set.
from hwk5_2.classifiers import Softmax
results = {}
best_val = -1
best_softmax = None
learning_rates = [1e-7, 5e-7]
regularization_strengths = [2.5e4, 5e4]

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength.
# This should be identical to the validation that you did for the SVM; save
# the best trained softmax classifier in best_softmax.
#####
for lr in learning_rates:
    for rs in regularization_strengths:
        softmax = Softmax()
        softmax.train(X_train, y_train, learning_rate = lr, reg=rs, num_iters = 1500,
            verbose = True)

```

```

y_pred_train = softmax.predict(X_train)
train_accuracy = np.mean(y_pred_train == y_train)

y_pred_val = softmax.predict(X_val)
current_accuracy = np.mean(y_pred_val == y_val)
results[(lr, rs)] = (train_accuracy, current_accuracy)

if current_accuracy > best_val:
    best_val = current_accuracy
    best_softmax = softmax
#####
#                               END OF YOUR CODE                               #
#####

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' % best_val)

iteration 0 / 1500: loss 389.585567
iteration 100 / 1500: loss 235.882251
iteration 200 / 1500: loss 143.151071
iteration 300 / 1500: loss 87.381301
iteration 400 / 1500: loss 53.580661
iteration 500 / 1500: loss 33.268365
iteration 600 / 1500: loss 20.905002
iteration 700 / 1500: loss 13.436617
iteration 800 / 1500: loss 8.994056
iteration 900 / 1500: loss 6.240335
iteration 1000 / 1500: loss 4.556667
iteration 1100 / 1500: loss 3.502107
iteration 1200 / 1500: loss 2.949472
iteration 1300 / 1500: loss 2.544311
iteration 1400 / 1500: loss 2.418960
iteration 0 / 1500: loss 779.154900
iteration 100 / 1500: loss 286.144912
iteration 200 / 1500: loss 105.892203
iteration 300 / 1500: loss 40.068351
iteration 400 / 1500: loss 16.085616
iteration 500 / 1500: loss 7.259269
iteration 600 / 1500: loss 3.952037
iteration 700 / 1500: loss 2.717868
iteration 800 / 1500: loss 2.311429
iteration 900 / 1500: loss 2.142549
iteration 1000 / 1500: loss 2.126397

```

```

iteration 1100 / 1500: loss 2.125953
iteration 1200 / 1500: loss 2.129377
iteration 1300 / 1500: loss 2.082927
iteration 1400 / 1500: loss 2.077699
iteration 0 / 1500: loss 392.871800
iteration 100 / 1500: loss 33.199343
iteration 200 / 1500: loss 4.451756
iteration 300 / 1500: loss 2.162057
iteration 400 / 1500: loss 2.076185
iteration 500 / 1500: loss 2.017442
iteration 600 / 1500: loss 2.011248
iteration 700 / 1500: loss 1.999447
iteration 800 / 1500: loss 2.119726
iteration 900 / 1500: loss 1.980435
iteration 1000 / 1500: loss 2.028366
iteration 1100 / 1500: loss 1.883927
iteration 1200 / 1500: loss 2.058723
iteration 1300 / 1500: loss 2.012429
iteration 1400 / 1500: loss 1.952310
iteration 0 / 1500: loss 767.466308
iteration 100 / 1500: loss 6.868844
iteration 200 / 1500: loss 2.115490
iteration 300 / 1500: loss 2.108528
iteration 400 / 1500: loss 2.123933
iteration 500 / 1500: loss 2.103379
iteration 600 / 1500: loss 2.134838
iteration 700 / 1500: loss 2.059736
iteration 800 / 1500: loss 2.072398
iteration 900 / 1500: loss 2.090008
iteration 1000 / 1500: loss 2.067229
iteration 1100 / 1500: loss 2.091932
iteration 1200 / 1500: loss 2.117153
iteration 1300 / 1500: loss 2.076470
iteration 1400 / 1500: loss 2.099181
lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.346367 val accuracy: 0.370000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.333041 val accuracy: 0.349000
lr 5.000000e-07 reg 2.500000e+04 train accuracy: 0.341429 val accuracy: 0.360000
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.330980 val accuracy: 0.353000
best validation accuracy achieved during cross-validation: 0.370000

```

```

In [23]: # evaluate on test set
         # Evaluate the best softmax on test set
         y_test_pred = best_softmax.predict(X_test)
         test_accuracy = np.mean(y_test == y_test_pred)
         print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))

softmax on raw pixels final test set accuracy: 0.356000

```

Inline Question - True or False

It's possible to add a new datapoint to a training set that would leave the SVM loss unchanged, but this is not the case with the Softmax classifier loss.

Your answer: True

Your explanation: In SVM, if the score for new datapoint is out of margin rangem the loss stays unchanged. However, in softmax, the loss will change.

```
In [24]: # Visualize the learned weights for each class
w = best_softmax.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)

w_min, w_max = np.min(w), np.max(w)

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])
```

